

Counterfactual Training

Teaching Models Plausible and Actionable Explanations

Patrick Altmeyer

Aleksander Buszydlík
Cynthia C. S. Liem

Arie van Deursen

Delft University of Technology

2026-03-23

Dynamic HTML Version

For the HTML version of these slides, visit patalt.org.

Training Opaque Models

Tweaking Parameters

Objective:

$$\min_{\theta} \{ \text{yloss}(M_{\theta}(\mathbf{x}), \mathbf{y}) \}$$

Training Opaque Models


Tweaking Parameters

Objective:

$$\min_{\theta} \{\text{yloss}(M_{\theta}(\mathbf{x}), \mathbf{y})\}$$

Solution:

$$\begin{aligned}\theta_{t+1} &= \theta_t - \nabla_{\theta} \{\text{yloss}(M_{\theta}(\mathbf{x}), \mathbf{y})\} \\ \theta^* &= \theta_T\end{aligned}$$




Explaining Opaque Models

Tweaking Inputs

Objective:

$$\min_{\mathbf{x}} \{ \text{yloss}(M_{\theta^*}(\mathbf{x}), \mathbf{y}^+) + \lambda \text{reg}(\mathbf{x}; \cdot) \}$$



Explaining Opaque Models

Tweaking Inputs


Objective:

$$\min_{\mathbf{x}} \{ \text{yloss}(M_{\theta^*}(\mathbf{x}), \mathbf{y}^+) + \lambda \text{reg}(\mathbf{x}; \cdot) \}$$

Solution:

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \nabla_{\mathbf{x}} \{ \text{yloss}(M_{\theta^*}(\mathbf{x}), \mathbf{y}^+) + \lambda \text{reg}(\mathbf{x}; \cdot) \}$$

$$\mathbf{x}^* = \mathbf{x}_T$$



Explanation or Adversarial Example?

Plausibility at all cost?

All of these counterfactuals are valid explanations for the model's prediction.

Pick your poison ...




Figure 1: Turning a 9 into a 7: Counterfactual explanations for an image classifier using different approaches (Altmeyer et al. 2024).

Faithful First, Plausible Second




Figure 2: Turning a 9 into a 7. *ECCCo* applied to MLP (a), Ensemble (b), JEM (c), JEM Ensemble (d).

- Insight:** faithfulness facilitates¹
- model quality checks (Figure 2).
 - state-of-the-art plausibility (Figure 3).




Figure 3: Results for different generators (from 3 to 5).

¹  Altmeyer et al. (2024) @ AAAI 2024. [blog]

Putting it all together

Counterfactual Training

First, *Tweaking Inputs*²

$$\begin{aligned}\mathbf{x}_{t+1} &= \mathbf{x}_t - \nabla_{\mathbf{x}} \{ECCCo(M_{\theta^*}(\mathbf{x}), \mathbf{y}^+)\} \\ \mathbf{x}^* &= \mathbf{x}_T\end{aligned}$$

Then, *Tweaking Parameters*

$$\begin{aligned}\theta_{t+1} &= \theta_t - \nabla_{\theta} \{\text{yloss}(M_{\theta}(\mathbf{x}), \mathbf{y}) + \text{div}(\mathbf{x}^*, \mathbf{x}^+, y^+; \theta)\} \\ \theta^* &= \theta_T\end{aligned}$$

²Generate faithful explanations using *ECCCo* objective (Altmeyer et al. 2024).

Counterfactual Training

- 1 Contrast faithful CE with data → **Explainability** ↑

Counterfactual Training

- 1 Contrast faithful CE with data → **Explainability** ↑
- 2 Feature mutability constraints → **Actionability** ↑ (holds provably under certain assumptions)

Counterfactual Training

- 1 Contrast faithful CE with data → **Explainability** ↑
- 2 Feature mutability constraints → **Actionability** ↑ (holds provably under certain assumptions)
- 3 Bonus: use nascent CE as AE → **Robustness** ↑

Counterfactual Training

- 1 Contrast faithful CE with data → **Explainability** ↑
- 2 Feature mutability constraints → **Actionability** ↑ (holds provably under certain assumptions)
- 3 Bonus: use nascent CE as AE → **Robustness** ↑

Counterfactual Training

- 1 Contrast faithful CE with data $\rightarrow$ **Explainability** $\uparrow$
- 2 Feature mutability constraints $\rightarrow$ **Actionability** $\uparrow$ (holds provably under certain assumptions)
- 3 Bonus: use nascent CE as AE $\rightarrow$ **Robustness** $\uparrow$




Figure 4: (a) conventional training, all mutable; (b) CT, all mutable; (c) conventional, *age* immutable; (d) CT, *age* immutable.

Counterfactual Training: Results



Figure 5: **Plausibility**: Visual explanations (counterfactuals) for baseline (top row) vs CT (bottom).




Figure 6: **Actionability**: Visual explanations (integrated gradients) as before. Five top and bottom rows immutable.




Figure 7: Test accuracies on adversarially perturbed data with varying perturbation sizes.

The Hard Numbers

Extensive experiments and ablation studies on nine datasets—synthetic, tabular and vision—generating millions of counterfactuals:³

- 1 **Plausibility** of CEs increases by up to 90%.
- 2 **Actionability**: cost of reaching valid counterfactuals with protected features decreases by 19% on average.
- 3 Models' adversarial **robustness** improves consistently.

³Facilitated by our CounterfactualExplanations.jl (Altmeyer et al. 2023) with multi-processing support and (DHPC) (2022).

Check it out!



Figure 8: Preprint



Figure 9: Software



Figure 10: Homepage

References

- Altmeyer, Patrick, Arie van Deursen, and Cynthia C. S. Liem. 2023. "Explaining Black-Box Models through Counterfactuals." *Proceedings of the JuliaCon Conferences* 1: 130.
<https://doi.org/10.21105/jcon.00130>.
- Altmeyer, Patrick, Mojtaba Farmanbar, Arie van Deursen, and Cynthia C. S. Liem. 2024. "Faithful Model Explanations through Energy-Constrained Conformal Counterfactuals." *Proceedings of the Thirty-Eighth AAAI Conference on Artificial Intelligence* 38: 10829–37.
<https://doi.org/10.1609/aaai.v38i10.28956>.
- (DHPC), Delft High Performance Computing Centre. 2022. *DelftBlue Supercomputer (Phase 1)*.
<https://www.tudelft.nl/dhpc/ark:/44463/DelftBluePhase1>.