ECCCos from the Black Box

Faithful Model Explanations through Energy-Constrained Conformal Counterfactuals

Patrick Altmeyer   Mojtaba Farmanbar   Arie van Deursen
Cynthia C. S. Liem

Delft University of Technology

2024-05-01

Faithfulness first, plausibility second.
OOOO

Reconciling Faithfulness and Plausibility
OOOO

Results
OOO

Questions?
OOOO

Pick your Poison

All of these counterfactuals are valid explanations for the model's prediction.

Which one would you pick?

Factual Wachter (p=0.9) Schut (p=0.92) REVISE (p=0.97)

Figure 1: Turning a 9 into a 7: Counterfactual explanations for an image classifier produced using *Wachter* (Wachter, Mittelstadt, and Russell 2017), *Schut* (Schut et al. 2021) and *REVISE* (Joshi et al. 2019).

Faithfulness first, plausibility second.

Faithfulness first, plausibility second.

We propose *ECCCo*: a new way to generate faithful model explanations that are as plausible as the underlying model permits.

Faithfulness first, plausibility second.
○○●○

Reconciling Faithfulness and Plausibility
○○○○

Results
○○○

Questions?
○○○○

Summary

- **Idea**: generate counterfactuals that are consistent with what the model has learned about the data.

Faithfulness first, plausibility second.
○○●○

Reconciling Faithfulness and Plausibility
○○○○

Results
○○○

Questions?
○○○○

Summary

- **Idea**: generate counterfactuals that are consistent with what the model has learned about the data.
- **Method**: constrain the model's energy and predictive uncertainty for the counterfactual.

Faithfulness first, plausibility second.
○○●○

Reconciling Faithfulness and Plausibility
○○○○

Results
○○○

Questions?
○○○○

Summary

- **Idea**: generate counterfactuals that are consistent with what the model has learned about the data.
- **Method**: constrain the model's energy and predictive uncertainty for the counterfactual.
- **Result**: faithful counterfactuals that are as plausible as the model permits.

Faithfulness first, plausibility second.
○○●○

Reconciling Faithfulness and Plausibility
○○○○

Results
○○○

Questions?
○○○○

Summary

- **Idea**: generate counterfactuals that are consistent with what the model has learned about the data.
- **Method**: constrain the model's energy and predictive uncertainty for the counterfactual.
- **Result**: faithful counterfactuals that are as plausible as the model permits.
- **Benefits**: enable us to distinguish trustworthy from unreliable models.

Faithfulness first, plausibility second.
○○○●

Reconciling Faithfulness and Plausibility
○○○○

Results
○○○

Questions?
○○○○

Counterfactual Explanations

$$\min_{\mathbf{Z}' \in \mathcal{Z}^L} \{\mathrm{yloss}(M_\theta(f(\mathbf{Z}')), \mathbf{y}^+) + \lambda \mathrm{cost}(f(\mathbf{Z}'))\}$$

Counterfactual Explanations (CE)
explain how inputs into a model
need to change for it to produce
different outputs.

Figure 2: Gradient-based
counterfactual search.

Reconciling Faithfulness and Plausibility

Faithfulness first, plausibility second.
0000

Reconciling Faithfulness and Plausibility
0●00

Results
000

Questions?
0000

Plausibility

Definition (Plausible Counterfactuals)

Let $\mathcal{X}|\mathbf{y}^+ = p(\mathbf{x}|\mathbf{y}^+)$ denote the true conditional distribution of samples in the target class $\mathbf{y}^+$. Then for $\mathbf{x}'$ to be considered a plausible counterfactual, we need: $\mathbf{x}' \sim \mathcal{X}|\mathbf{y}^+$.

Why Plausibility?

Plausibility is positively associated with actionability, robustness (Artelt et al. 2021) and causal validity (Mahajan, Tan, and Sharma 2020).


Figure 3: Kernel density estimate (KDE) for the conditional distribution, $p(\mathbf{x}|\mathbf{y}^+)$, based on observed data. Counterfactual path as in Figure 2.

Faithfulness first, plausibility second.
○○○○

Reconciling Faithfulness and Plausibility
○○●○

Results
○○○

Questions?
○○○○

Faithfulness

Definition (Faithful Counterfactuals)

Let $\mathcal{X}_\theta | \mathbf{y}^+ = p_\theta(\mathbf{x}|\mathbf{y}^+)$ denote the conditional distribution of $\mathbf{x}$ in the target class $\mathbf{y}^+$, where θ denotes the parameters of model M_θ. Then for $\mathbf{x}'$ to be considered a faithful counterfactual, we need:
$\mathbf{x}' \sim \mathcal{X}_\theta | \mathbf{y}^+$.

Trustworthy Models

If the model posterior approximates the true posterior $\left(p_\theta(\mathbf{x}|\mathbf{y}^+) \to p(\mathbf{x}|\mathbf{y}^+)\right)$, faithful counterfactuals are also plausible.


Figure 4: KDE for learned conditional distribution, $p_\theta(\mathbf{x}|\mathbf{y}^+)$. Yellow stars indicate conditional samples generated through SGLD for a joint energy model (JEM).

Faithfulness first, plausibility second.
○○○○

Reconciling Faithfulness and Plausibility
○○○●

Results
○○○

Questions?
○○○○

ECCCo

Key Idea

Use the hybrid objective of joint energy models (JEM) and a model-agnostic penalty for predictive uncertainty: Energy-Constrained ($\mathcal{E}_\theta$) Conformal (Ω) Counterfactuals (*ECCCo*).

ECCCo objective[a]:

$$\min_{\mathbf{Z}' \in \mathcal{Z}^L} \{ L_{\text{clf}}(f(\mathbf{Z}'); M_\theta, \mathbf{y}^+) + \lambda_1 \text{cost}(f(\mathbf{Z}'))$$
$$+ \lambda_2 \mathcal{E}_\theta(f(\mathbf{Z}')|\mathbf{y}^+) + \lambda_3 \Omega(C_\theta(f(\mathbf{Z}'); \alpha)) \}$$


Figure 5: Gradient fields and counterfactual paths for different generators.

---

[a]We leverage ideas from Grathwohl et al. (2020) and Stutz et al. (2022). See the paper and appendix for a derivation of the objective from first principles.

Faithfulness first, plausibility second.
OOOO

Reconciling Faithfulness and Plausibility
OOOO

Results
●OO

Questions?
OOOO

# Results

Visual Evidence

(a)   (b)

(c)   (d)

Figure 6: Turning a 9 into a 7. *ECCCo* applied to MLP (a), Ensemble (b), JEM (c), JEM Ensemble (d).

ECCCo generates counterfactuals that

- faithfully represent model quality (Figure 6).
- achieve state-of-the-art plausibility (Figure 7).

Factual   ECCCo   ECCCo+   REVISE   Schut   Wachter

Figure 7: Results for different generators (from 3 to 5).

Faithfulness first, plausibility second.
○○○○

Reconciling Faithfulness and Plausibility
○○○○

Results
○○●

Questions?
○○○○

The Numbers

- Large benchmarks on a variety of models and datasets from various domains.
- *ECCCo* achieves state-of-the-art faithfulness across models and datasets and approaches state-of-the-art plausibility for more trustworthy models.

Model	Generator	California Housing			GMSC		
		Unfaithfulness ↓	Implausibility ↓	Uncertainty ↓	Unfaithfulness ↓	Implausibility ↓	Uncertainty ↓
MLP Ensemble	ECCCo	**3.69 ± 0.08****	1.94 ± 0.13	**0.09 ± 0.01****	3.84 ± 0.07**	2.13 ± 0.08	**0.23 ± 0.01****
	ECCCo+	3.88 ± 0.07**	1.20 ± 0.09	0.15 ± 0.02	**3.79 ± 0.05****	1.81 ± 0.05	0.30 ± 0.01*
	ECCCo (no CP)	3.70 ± 0.08**	1.94 ± 0.13	0.10 ± 0.01**	3.85 ± 0.07**	2.13 ± 0.08	0.23 ± 0.01**
	ECCCo (no EBM)	4.03 ± 0.07	1.12 ± 0.12	0.14 ± 0.01**	4.08 ± 0.06	0.97 ± 0.08	0.31 ± 0.01*
	REVISE	3.96 ± 0.07*	**0.58 ± 0.03****	0.17 ± 0.03	4.09 ± 0.07	**0.63 ± 0.02****	0.33 ± 0.06
	Schut	4.00 ± 0.06	1.15 ± 0.12	0.10 ± 0.01**	4.04 ± 0.08	1.21 ± 0.08	0.30 ± 0.01*
	Wachter	4.04 ± 0.07	1.13 ± 0.12	0.16 ± 0.01	4.10 ± 0.07	0.95 ± 0.08	0.32 ± 0.01
JEM Ensemble	ECCCo	1.40 ± 0.08**	0.69 ± 0.05**	0.11 ± 0.00**	1.20 ± 0.06*	0.78 ± 0.07**	0.38 ± 0.01
	ECCCo+	**1.28 ± 0.08****	0.60 ± 0.04**	0.11 ± 0.00**	**1.01 ± 0.07****	0.70 ± 0.07**	0.37 ± 0.01
	ECCCo (no CP)	1.39 ± 0.08**	0.69 ± 0.05**	0.11 ± 0.00**	1.21 ± 0.07*	0.77 ± 0.07**	0.39 ± 0.01
	ECCCo (no EBM)	1.70 ± 0.09	0.99 ± 0.08	0.14 ± 0.00*	1.31 ± 0.07	0.97 ± 0.10	0.32 ± 0.01**
	REVISE	1.39 ± 0.15**	**0.59 ± 0.04****	0.25 ± 0.07	1.01 ± 0.07**	**0.63 ± 0.04****	0.33 ± 0.07
	Schut	1.59 ± 0.10*	1.10 ± 0.06	**0.09 ± 0.00****	1.34 ± 0.07	1.21 ± 0.10	**0.26 ± 0.01****
	Wachter	1.71 ± 0.09	0.99 ± 0.08	0.14 ± 0.00	1.31 ± 0.08	0.95 ± 0.10	0.33 ± 0.01

Table 1: Results for tabular datasets: sample averages +/- one standard deviation across valid counterfactuals. The best outcomes are highlighted in bold. Asterisks indicate that the given value is more than one (*) or two (**) standard deviations away from the baseline (*Wachter*).

Questions?

Faithfulness first, plausibility second.
OOOO

Reconciling Faithfulness and Plausibility
OOOO

Results
OOO

Questions?
O●OO

Questions?

With thanks to my co-authors Mojtaba Farmanbar, Arie van Deursen and
Cynthia C. S. Liem.

Code

The code used to run the analysis for this work is built on top of `CounterfactualExplanations.jl`.

There is also a corresponding paper, *Explaining Black-Box Models through Counterfactuals*, which has been published in JuliaCon Proceedings.

Figure 8: Trustworthy AI in Julia: github.com/JuliaTrustworthyAI

References

Artelt, André, Valerie Vaquet, Riza Velioglu, Fabian Hinder, Johannes Brinkrolf, Malte Schilling, and Barbara Hammer. 2021. "Evaluating Robustness of Counterfactual Explanations." In *2021 IEEE Symposium Series on Computational Intelligence (SSCI)*, 01–09. IEEE.

Grathwohl, Will, Kuan-Chieh Wang, Joern-Henrik Jacobsen, David Duvenaud, Mohammad Norouzi, and Kevin Swersky. 2020. "Your Classifier Is Secretly an Energy Based Model and You Should Treat It Like One." In *International Conference on Learning Representations*.

Joshi, Shalmali, Oluwasanmi Koyejo, Warut Vijitbenjaronk, Been Kim, and Joydeep Ghosh. 2019. "Towards Realistic Individual Recourse and Actionable Explanations in Black-Box Decision Making Systems." https://arxiv.org/abs/1907.09615.

Mahajan, Divyat, Chenhao Tan, and Amit Sharma. 2020. "Preserving Causal Constraints in Counterfactual Explanations for Machine